

Multitype Messages in the Same Kafka Topic

Options, Challenges, and Techniques, Apr 2021

Agenda

1. Prelude
2. Why Multitype in the Same Topic
3. Options
4. Challenges
5. Techniques

Prelude

- Kafka, Schema Registry, Avro, Java
- We compile Avro schemas into Java classes - `SpecificRecord`
- One-topic-one-type
 - just like in database, one-table-one-entity

👉 In this slide, we use `SpecificRecord` to refer to those generated model classes

Why Multitype in the Same Topic

- Business nature
 - Events are naturally multityped within business domains
 - Events evolve as business evolve
- **Ordering !!!**
 - Ordering is CRITICAL to event based processing
 - have-your-cake \Rightarrow eat-it \Rightarrow 🧑
 - eat-it \Rightarrow have-your-cake \Rightarrow 🎉
 - Kafka reserves the order of events ONLY if they are in the same partition
 - Being in the same topic is a prerequisite for being in the same partition

👉 There are other prerequisites for guaranteeing ordering in Kafka

Options

- JSON
- Tuple of KVP
- Map
- Union
- Subject Naming Strategy

Schemaless



Semi-schema



Schemaful

JSON

- 🧑 Zero Schema
- 👍 Super flexible
- 👎 Type is implicit
- 🤞 Good luck on producers & consumers
- This is effectively similar to using Compatibility Type NONE with Avro

👎 This is not saying one can't enforce schema with JSON

Tuple & Map

- 🧑 One Schema for All
- 👍 Unique keys guarantee (Map)
- 🙅 Potentially non-unique keys (Tuple)
- 🧑 Challenge supporting various value types
- 🏆 Tricky on nested structure (especially for Tuple)
- 🙅 Type is implicit
- 🙅 Responsibilities on producers & consumers

🙅 GraphQL does NOT support Map

Union

- In computer science, a union is a value that may have any of several representations or formats within the same position in memory. — Wikipedia
- In Avro, a union indicates that a field might have more than one data type. E.g.

```
{
  "type" : "record",
  "name" : "Candidate",
  "fields" :
  [
    {
      "name" : "experience",
      "type": ["null", "int"]
    }
  ]
}
```


Union for Multitype

In below schema, payload holds either a Sms or Email type

```
{
  "type": "record",
  "name": "Event",
  "namespace": "demo.model",
  "fields": [
    {
      "name": "eventId", "type": "string"
    },
    {
      "name": "payload",
      "type": [
        {
          "type": "record",
          "name": "Sms",
          "fields": [
            {
              "name": "phoneNumber", "type": "string"
            },
            {
              "name": "text", "type": "string"
            }
          ]
        },
        {
          "type": "record",
          "name": "Email",
          "fields": [
            {
              "name": "address", "type": "string"
            },
            {
              "name": "title", "type": "string"
            }
          ]
        }
      ]
    }
  ]
}
```

Subject Naming Strategy

- What is **Subject**?
 - **Q** : Can I use schema-registry without Kafka broker independently? (Schema Registry Issue #533)
 - **A** : Definitely. That's why we use slightly different terminology in the schema registry ("subjects") than we use in Kafka ("topics").
- Often people want one-topic-one-type, meaning one-topic-one-subject
 - Hence the default strategy is: `TopicNameStrategy`
- Two other strategies
 - `RecordNameStrategy`
 - `TopicRecordNameStrategy`
 - 👍 Types evolve independently

Naming Strategies in Action

- Producer publishes customer activities to topic activity
 - TopicNameStrategy (default)
 - activity-value, only one type of message allowed
 - RecordNameStrategy
 - demo.model.MonetaryActivity
 - demo.model.NonMonetaryActivity
 - ...
 - TopicRecordNameStrategy
 - activity-demo.model.MonetaryActivity
 - activity-demo.model.NonMonetaryActivity
 - ...

👉 Control Center UI supports TopicNameStrategy only

Naming Strategy in Code

- Producer

```
props.put("value.subject.name.strategy", TopicRecordNameStrategy.class);
KafkaProducer producer = new KafkaProducer(props);

// Subject: activity-demo.model.MonetaryActivity
producer.send(new ProducerRecord("activity", "key1", new MonetaryActivity()));



// Subject: activity-demo.model.NonMonetaryActivity (same topic as above)
producer.send(new ProducerRecord("activity", "key2", new NonMonetaryActivity()));
```

- Consumer

```
ConsumerRecords records = consumer.poll(ONE_SECOND);
for (ConsumerRecord record : records) {
    Object activity = record.value();
    if (activity instanceof MonetaryActivity) { ... }
    if (activity instanceof NonMonetaryActivity) { ... }
}
```

👉 Generics omitted for brevity

Challenges in Union and Subject Naming Strategy

-  Message Filtering
 - Similar to `SELECT...WHERE...` in SQL, and Message Selector in JMS
 - There is NO BUILT-IN way to filter out unwanted messages
- Add a new Type
 - Breaks existing consumers if
 - Using `SpecificRecord`
 - AND filtering not implemented
- Deprecate a Type
 - Breaks new consumers for the same reason
-  Evolve types inside Union

Message Filtering

- Because Kafka brokers are 'DUMB'
- and Kafka requires that a consumer to take ALL messages it receives
 - Consumer's Life:
 1. fetch messages in bytes
 2. deserialize into objects
 3. process them
 4. commit offset
 5. goto step 1
 - 😞 Consumer cannot say no to unwanted bytes
- 🙅 Consequently, Introducing new types BREAKS existing consumers using `SpecificRecord`
 - Unless the consumer upgrades to newer schema or implements custom filtering

👉 This applies to both the Union and Subject Naming Strategy approaches

Techniques - GenericRecord

- Consumer use `GenericRecord` rather than `SpecificRecord`
- This is quite similar to Map approach
- 👍 Deserialisation always SUCCEEDS

👉 One can have producers publishing messages of new type using `SpecificRecord` while not breaking existing consumers from consuming them in `GenericRecord`

Kafka Headers

- **Header** exists in many places, TCP, HTTP, HTML, JMS, SOAP, etc.
- Headers are **Metadata**, serve purposes like routing, authentication, etc.
- Header support added to Kafka in version 0.11 in 2017

```
public interface Header {  
    String key();  
    byte[] value();  
}
```

- Available to Producer and Consumer

```
record.headers().add("meta1", "important-data".getBytes(UTF_8)); // Producer  
  
for (Header header : record.headers()) { ... } // Consumer
```

- Header support added to Deserialiser interface in 2018

```
public interface Deserializer<T> {  
    void configure(Map<String, ?> configs, boolean isKey);  
    T deserialize(String topic, byte[] data);  
    T deserialize(String topic, Headers headers, byte[] data);  
}
```


Techniques - Swallow But Not Digest with Header Filtering

- 🤝 Filter out unwanted messages
 - Producer tags a message's type via header

```
record.headers().add("Type", "MonetaryActivity".getBytes());  
producer.send(record);
```

- Deserializer converts them into null as per tag value

```
KafkaAvroDeserializer worker;  
  
public Object deserialize(String topic, Headers headers, byte[] data) {  
    if (!isKnown(headers)) return null;  
  
    return worker.deserialize(topic, headers, data);  
}
```

- Consumer ignores null messages and continue

```
props.put("value.deserializer", TolerantDeserializer.class);  
props.put("tolerant.headerName", "Type");  
props.put("tolerant.headerValueRegex", "MonetaryActivity|NonMonetaryActivity");  
  
for (ConsumerRecord<String, SpecificRecordBase> record : records) {  
    if (Objects.isNull(record.value())) continue;  
}
```

👉 This technique can be applied to both Subject Naming Strategy and Union approaches

Safely Add New Types with Defensive Consumers

Producers	Topic	Consumers
ProducerA: E1	=> multi-events-topic =>	ConsumerA: E1 🌧️
ProducerB: E2		ConsumerB: E1 E2 🌧️
ProducerC: E-new		ConsumerC: E1 E-new

Conclusion

- It is Complicated!
- It Depends.